

FOUNDATIONAL AI COMPUTING · ALGEBRAIC TOPOLOGY

# Beyond the Transformer: The Native Geometric Foundation

Autoregressive geometric state models with exact addressed memory and multiplier-free serving kernels. Engineered from fundamental non-Euclidean manifolds for radical energy and thermal efficiency.

FOUNDER & RESEARCH LEAD

**Casey Allard**

casey.allard@r4labs.org

2+ Years Self-Funded Engineering in Rust

LABORATORY BASE

**R4 Laboratories Inc.**

Providence, RI · Distributed Nodes

Official Portal: <https://r4labs.org>

OPEN SOURCE CRATES

**Core Codebases**

Rust Engine: [github.com/R4-Laboratories](https://github.com/R4-Laboratories)

UOR Specifications & Topological Routers

THE COMPUTE & MEMORY CRISIS

# The Scaling Ceiling of Dense Transformers

Brute-force matrix scaling is colliding with the physical limits of memory bandwidth, power grids, and thermal dissipation.

BOTTLENECK 01

## Quadratic Memory Wall

Standard attention scales as  $O(N^2)$ . Sprawling KV-caches starve memory bandwidth, rendering large-context real-time inference economically and physically prohibitive.

BOTTLENECK 02

## Thermal & Energy Tax

Datacenters are power-constrained by megawatt allocations; edge devices overheat. Continuous floating-point matrix multiplication dissipates immense energy for incremental gains.

BOTTLENECK 03

## Probabilistic Hallucination

Unconstrained continuous vector approximations blur historical facts over long reasoning chains. Systems lack immutable, deterministic coordinate addressing for exact retrieval.

FIRST-PRINCIPLES COMPUTATION

# Computing via Topology, Symmetry & Exact Invariants

Intelligence does not require brute-force tensor multiplication. Nature and physics conserve energy through symmetry groups, topological invariants, and discrete transitions.

### Conventional Transformers (Status Quo)

- **Dense Parameter Soup:** Millions of unconstrained continuous weights requiring continuous optimization.
- **Continuous Floating-Point Multiplication:** Tensor multipliers run at every step, creating thermal saturation.
- **Runaway Cache Scaling:** Memory bandwidth choked by sprawling KV-cache growth.

### The R4 Geometric Paradigm

- **Structured Manifold Evolution:** States transition over 4D hyperspheres ( $S^3$ ) and paired golden icosians ( $H4 \oplus \phi H4$ ).
- **Multiplier-Free Serving:** Served kernels execute strictly via integer add/sub, bitwise shifts, and table lookups.
- **Exact Addressed Memory (UOR):** Deterministic coordinate lookup eliminates hallucination and memory bloat.

EXECUTION LIFECYCLE

# The Four Fundamental Primitives

From token ingestion to served response: a complete ground-up reimagination of model architecture.

PRIMITIVE 01

### Prime Ingestion

Tokens mapped to exact prime addresses and fixed zeta-zero phases, preserving structural and harmonic identity without continuous vector dilution.

PRIMITIVE 02

### R4/S3 Manifold

Autoregressive state transitions over 4D hyperspheres and paired icosian lattices, preserving chiral orientation and exact algebraic invariants.

PRIMITIVE 03

### Exact Memory

Universal Object Reference (UOR) addressing. Information is anchored as deterministic coordinates, ending lossy associative attention decay.

PRIMITIVE 04

### Multiplier-Free

Served inference executes strictly via bounded integer addition, bitwise shifts, and compact table lookups. Zero floating-point multiplier instructions.

MEASURED ARCHITECTURAL ADVANTAGES

# Radical Efficiency by Mathematical Design

Eliminating continuous matrix multiplication unlocks unprecedented latency, memory, and power characteristics.

0

**Served Multipliers**  
Zero floating-point or transcendental multiplier instructions in served execution kernels.

10x+

**Energy Reduction**  
Sub-watt-hour token generation on consumer and edge hardware with near-zero thermal envelope.

0(1)

**Constant Memory**  
Bounded geometric state space replaces sprawling multi-gigabyte KV-cache memory footprints.

100%

**Deterministic**  
Bitwise deterministic artifact generation, exact addressable facts, and zero probabilistic decay.

ACCELERATING DISCRETE GEOMETRIC PRIMITIVES

# Why NVIDIA? Hardware Synergy & Optimization

Unlocking extreme throughput and density on NVIDIA Ada Lovelace, Hopper, and Blackwell architectures.

TENSOR CORES

## INT4 / INT8 Integer Pipelines

R4 serving kernels use bounded  $\leq 4$ -bit integer linear maps and table reads—mapping natively to high-density integer pipelines on NVIDIA Tensor Cores for massive throughput multiplication.

MEMORY BUS

## Memory Bandwidth Liberation

By eradicating sprawling KV-caches in favor of compact geometric addressed coordinates, R4 relieves High Bandwidth Memory (HBM) bottlenecks, supporting 10x more concurrent streams per GPU.

CUDA ACCELERATION

## Custom CUDA Ring Kernels

Opportunity to co-develop CUDA and TensorRT acceleration for golden ratio  $Z[\varphi]$  ring arithmetic, Hopf fibration projections, and prime coordinate routing directly on NVIDIA SMs.

OPEN SOURCE & PRODUCTION CODEBASE

# 2+ Years of Self-Funded Engineering in Rust

A robust, memory-safe codebase, mathematical formalizations, and active repositories under @R4-Laboratories.

CORE ENGINE

uor-r4

The primary experimental autoregressive geometric state model, Rust autodiff training tools, and multiplier-free serving kernels with bitwise deterministic artifacts.

MANIFOLD ROUTER

prime-router

Angular-manifold and prime-coordinate routing framework for high-dimensional semantic spaces and non-Euclidean state updates.

ADDRESSING SPEC

uor-addr-1 & Web Runtimes

Canonical content addressing standard (RFC8785) for agent-produced content, alongside in-browser WebAssembly native inference engines (uor-r4-wasm-chat).

CONTINUOUS SCIENTIFIC VERIFICATION

# Autonomous Multi-Agent Research Ecosystem

R4 Laboratories operates autonomous multi-agent research benches that enforce cryptographic receipts and eliminate reproducibility crises.

AUDITABILITY

### Cryptographic Evidence

Every hypothesis, training checkpoint, evaluation score, and parameter update is recorded in immutable, hash-linked receipt ledgers.

REPRODUCIBILITY

### Zero Verification Debt

Pinning exact compiler states, seeds, and training ladders guarantees bitwise identical results across diverse developer and cluster environments.

CONTINUOUS DISCOVERY

### Autonomous Bench Operations

Multi-agent benches coordinate across isolated worktrees, continuously exploring non-Euclidean topologies and hard/relaxed discretization bridges.

MARKET VERTICALS & REVENUE PATH

# Target Markets & Commercialization

Targeting high-margin, mission-critical sectors where transformers are physically non-viable or economically prohibitive.

VERTICAL 01

### Edge AI, Defense & Aerospace

Thermal-constrained and battery-powered platforms (drones, satellites, robotics, mobile) requiring local offline frontier reasoning without burning out batteries or overheating.

VERTICAL 02

### Enterprise Exact Memory

Financial audit, legal compliance, and healthcare systems requiring verifiable, non-hallucinatory retrieval of historical facts and regulatory documents with cryptographic proof.

VERTICAL 03

### Foundation Model IP Licensing

Licensing multiplier-free geometric model weights, topological routers, and ultra-lightweight serving runtimes to hyperscalers, chipmakers, and device manufacturers.

DEVELOPMENT TIMELINE

# Roadmap: From Theory to Frontier Scaling

A disciplined execution trajectory moving from Rust mathematical foundations to hardware-accelerated enterprise models.

Q4 2026 (CURRENT)

**Training Ladder**

Finalize Rust autodiff joint-learning ladder, hard/relaxed discretization bridge, and public open-source benchmark suite.

Q1 2027

**CUDA Acceleration**

Custom NVIDIA Tensor Core INT4/INT8 kernels; initial 1B-equivalent geometric model release with multiplier-free serving.

Q2 2027

**Edge Deployment**

Ultra-low power edge deployment on embedded silicon and consumer hardware; enterprise pilot trials for exact addressed memory.

Q3 2027

**Frontier Scaling**

Multi-node autonomous scaling and frontier capability evaluation across complex multi-step reasoning and mathematical benchmarks.

PARTNERSHIP OBJECTIVES

# The Ask: Accelerating R4 with NVIDIA Inception

Partnering with NVIDIA to scale native geometric computing on accelerated enterprise hardware.

REQUEST 01

### GPU Cluster Compute

Access to NVIDIA DGX Cloud or partner GPU resources (Lambda, CoreWeave, RunPod) to scale our offline Rust training ladder, parameter exploration, and empirical validation.

REQUEST 02

### Architectural Collaboration

Direct engagement with NVIDIA DevRel and engineers to optimize custom integer and table-lookup CUDA kernels for TensorRT and Blackwell integer pipelines.

REQUEST 03

### Inception Go-to-Market

Visibility through the NVIDIA Inception showcase, co-authored technical publications, and venture connections via Inception Capital Connect.

THE FUTURE OF COMPUTE

# The Geometric Future of AI

*"The future of intelligence will not be decided by who burns the most power, but by who discovers the most elegant geometry."*

PRINCIPAL RESEARCHER & FOUNDER

**Casey Allard**

casey.allard@r4labs.org

LABORATORY HEADQUARTERS

**Providence, RI · USA**

<https://r4labs.org>